

Inference throughput on RTX 5070 (bf16-mixed, batch=1, 100 warmup / 1000 measurement)

